

Real-time response & Real-time operating systems (RTOS)

Luis Gonzalez, PhD
Computer Science Department
Tecnologico de Monterrey, Campus Guadalajara

Outline

- Introduction
- Concurrency
- Real-time planning algorithms
- Reentrant code
- Design patterns

Outline

- **Introduction**
- Concurrency
- Real-time planning algorithms
- Reentrant code
- Design patterns

Introduction

- It is time to consider the issues involved in the accurate measurement of time
- These issues are important cause many embedded systems must satisfy real-time constraints
- Real-time applications imply **deterministic processing**, i.e., guaranteeing that a particular activity/task will always be completed in a **well-defined time interval**
- Most embedded SW is multitasking in nature, supporting multiple tasks running concurrently on the same CPU
- But one CPU can only run one instruction at a time
- But multiple tasks need to run simultaneously in an embedded application

Introduction

- Example : Temperature controller

- Components

- Temperature sensor
 - Control panel with keypad and LCD

- Specs

- The controller must maintain a room within a temperature range
 - If the temperature goes above that range, a red LED is turned on
 - If the temperature goes below that range, a green LED is turned on
 - Both LEDs are off if the temperature is within the temperature range
 - The current temperature is always displayed on the LCD
 - Users can reset the upper and lower limits of the temperature range at any time using the keypad



Introduction

- Specs suggest that the software must run at least two tasks concurrently

Task 1

```
while (1)
{
    sample ADC with sensor;
    rest temp;
    if (temp > up_T) turn on Red;
    else if (temp < low_T) turn on Green;
    else turn off both
    wait (100 ms);
}
```

Task 2

```
while (1)
{
    switch(key)
    {
        Up: Set up_T;
        break;
        Low: Set low_T;
        break;
    }
}
```

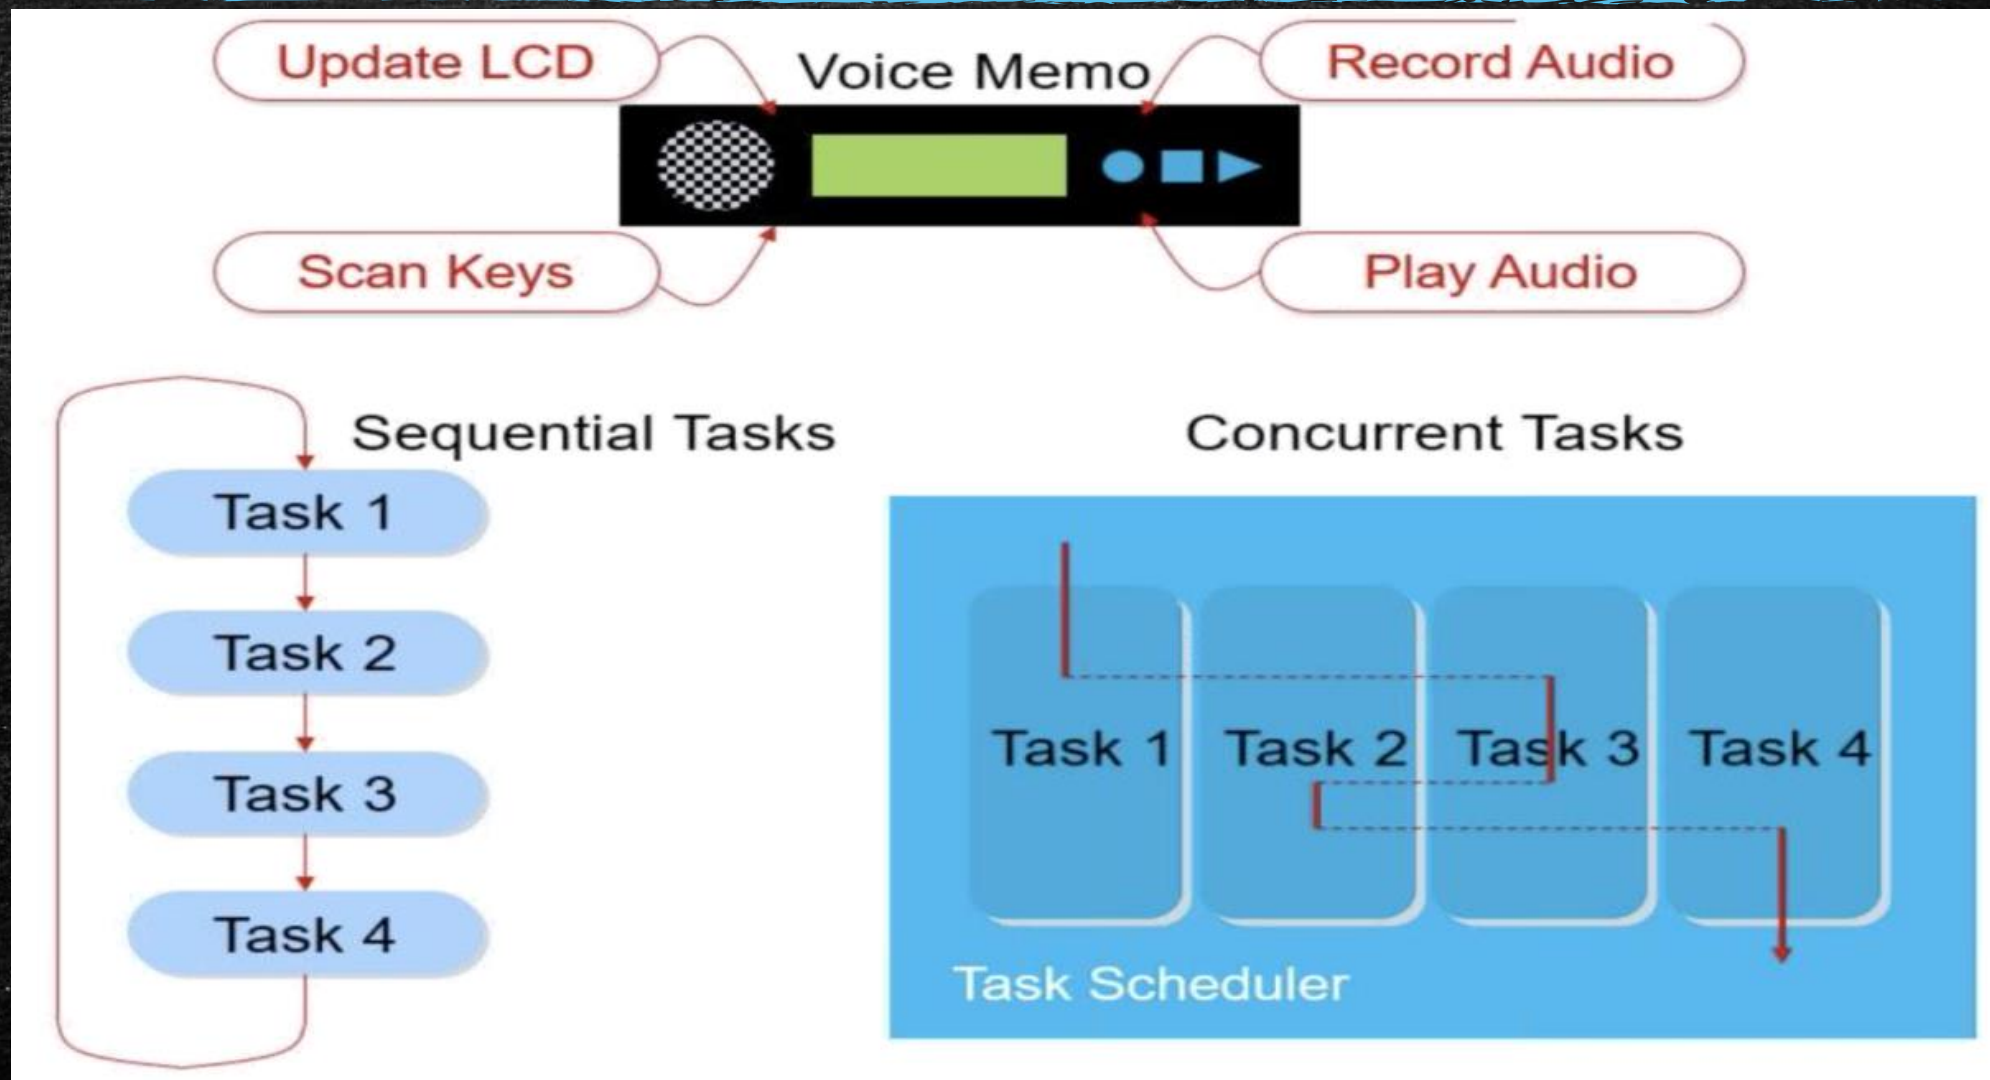

Introduction

- Task 1
 - monitors the current temperature and takes action when the temperature goes beyond the temperature range
 - Samples the temperature sensor every 100 ms
- Task 2
 - Take user inputs from the keypad to reset the temperature upper and lower limits

This very simple embedded system deals with real-time and multiple tasks. Two options can be identified to implement this system

- Real-Time operating system (RTOS)
- Using the interrupt mechanism supported by all microprocessors/microcontrollers

Why multitasking?



OS, RTOS or Baremetal?

- An **OS** provides a set of system and platform execution services for the app developer, specially around the management and use of target system resources, such as
 - Memory
 - Concurrency units (processes, tasks or threads)
 - Event queues
 - Interrupts
 - Hardware
 - Application programs

OS, RTOS or Baremetal?

- Most OSs do not provide any guarantee about timeliness
- Desktop OSs may invoke unpredictable delays due to internal processing, memory management or garbage collection at unforeseeable times
- This unpredictability makes them unsuitable for real-time and embedded environments that are constrained in time and resources

OS, RTOS or Baremetal?

- A real-time operating system (RTOS) is a multitasking operating system intended for real-time and embedded applications
- It is a piece of software that allows multiple tasks to **appear to run simulatenously**
- It also allows resources to be managed and shared
- They are written to provide services with good efficiency and performance
- Usually the **predictability of the performance** is more important than the maximum throughput

OS, RTOS or Baremetal?

- RTOSs are smaller and often less capable than desktop OSs
- RTOSs do not guarantee real-time performance but they provide an application environment so that appropriate developed applications can achieve real-time performance
- RTOSs run applications and tasks using one of three basic schemes
 - **Even-driven systems**
 - They handle events as they arise and schedule tasks to handle the processing
 - They use task priority as a quantitative means to determine which task will run if multiple tasks are ready to run
 - Task priorities are typically **static**, specified at design time
 - Task priorities can also be **dynamic**, varying the task priorities to account for the current operating conditions

OS, RTOS or Baremetal?

- RTOSs run applications and tasks using one of three basic schemes
 - Periodic runtime tasks (time-base schemes)
 - Cyclical task execution (sequence-based schemes)
- When the embedded system has too few resources to support an RTOS, the alternative is **baremetal**
- The application code must cope with managing services and functionality
- The application may simply be a set of interrupt handlers that communicate via a shared resource scheme such as queueing of shared memory

Outline

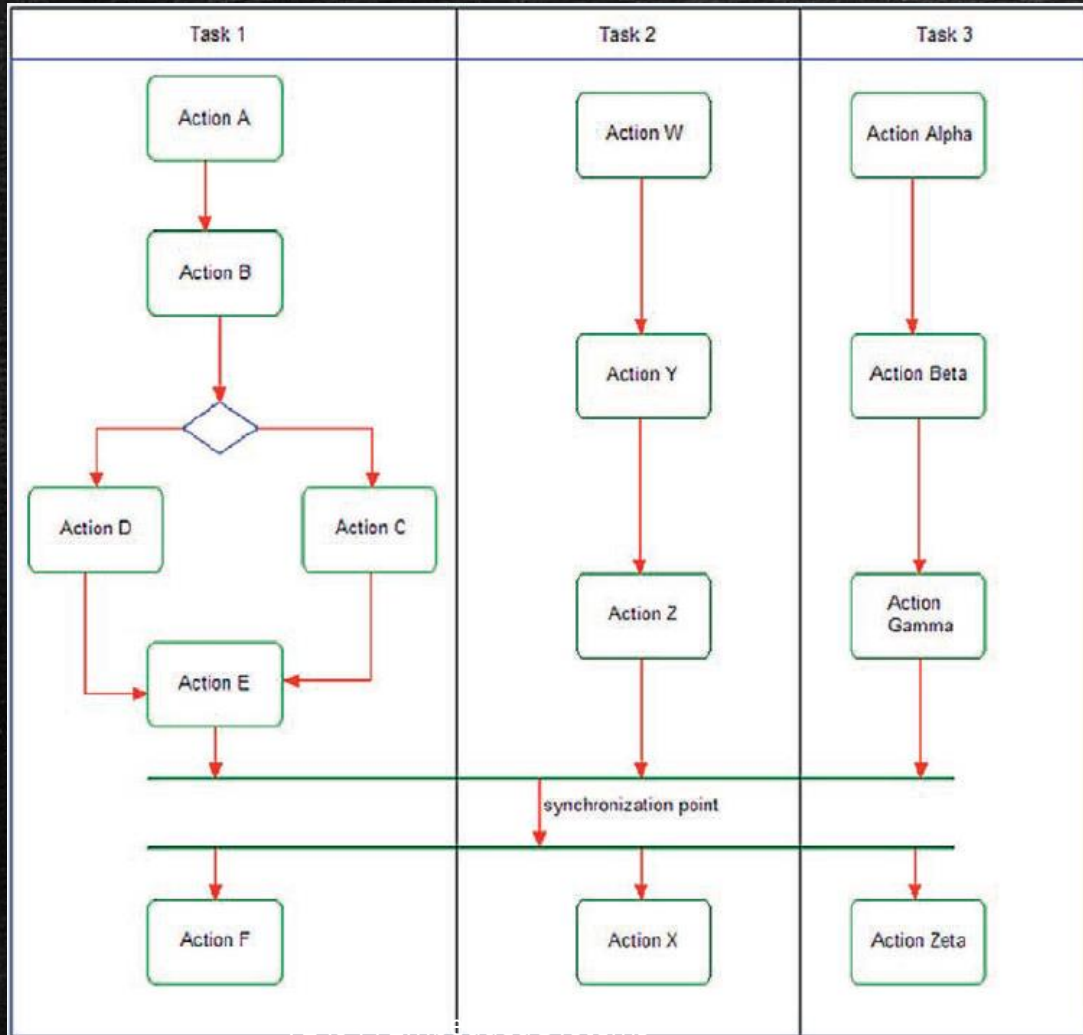
- Introduction
- Concurrency
- Real-time planning algorithms
- Reentrant code
- Design patterns

Concurrency

- Most embedded systems need to perform several activities **simultaneously**
- The key to achieve this is through the definition, understanding and management of the **concurrency model** of the system
- According to Wikipedia

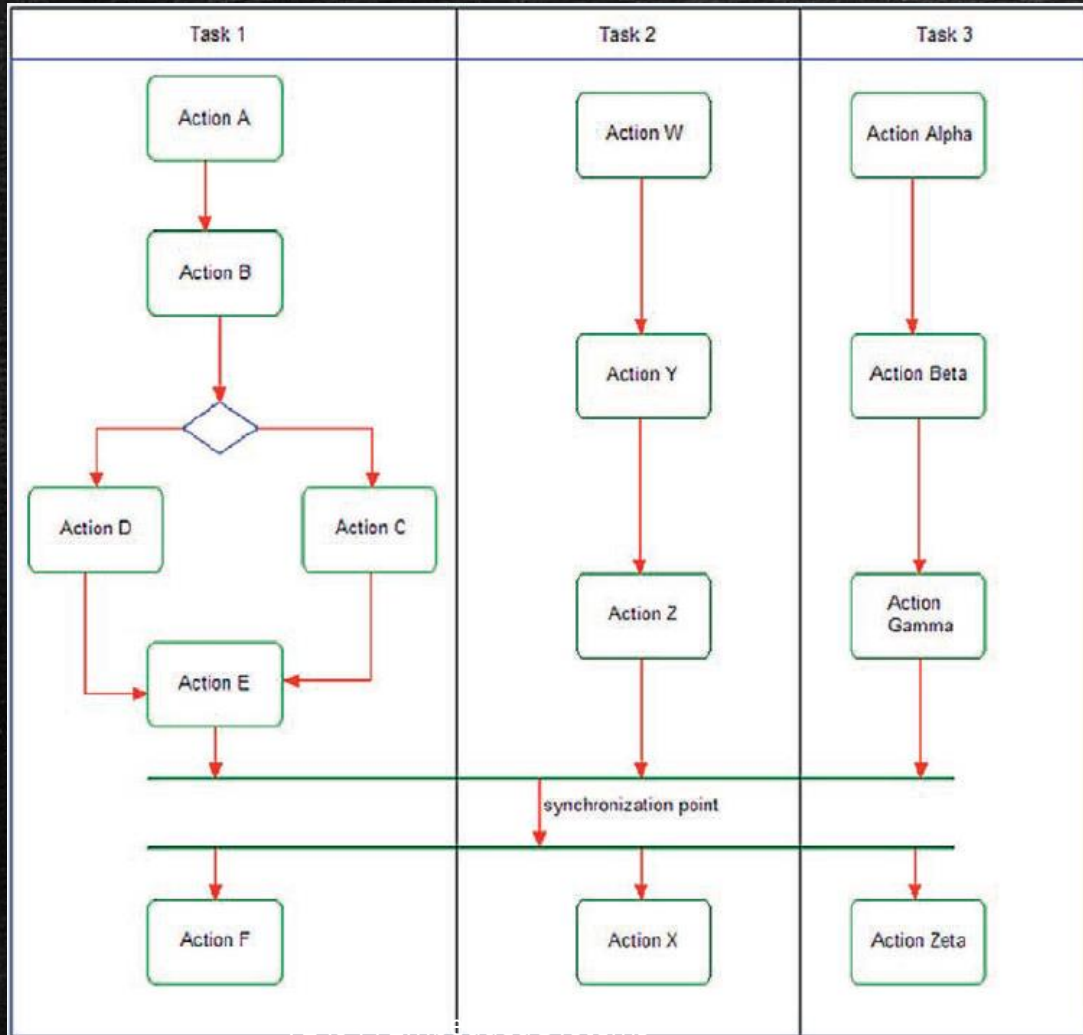
“Concurrency is the ability of different parts or units of a program, algorithm or problem to be executed out-of-order or in partial order, without affecting the final outcome. This allows for parallel execution of the **concurrent units**, which can significantly improve the overall speed of the execution in multi-processor and multi-core systems.”

Concurrency: Actions, action sequences and tasks



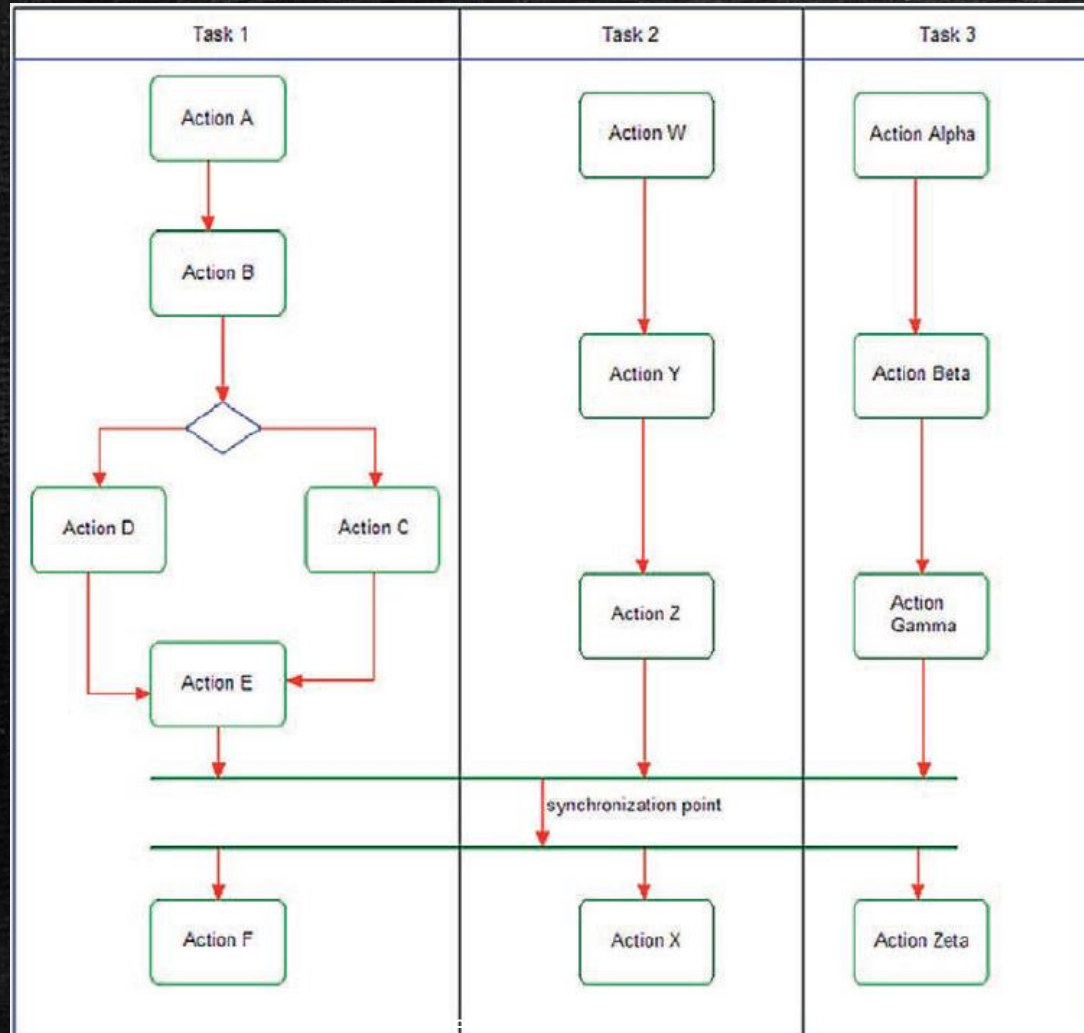
- A task (aka **concurrency unit**) consists of a set of actions performed in a specific sequence
- These **action sequences** are executed independently from the others
- Within a task, action sequences are fully known
 - Action A -> action B -> action C or action D -> ...

Concurrency: Actions, action sequences and tasks



- What is not known is the order of action execution **between** tasks
- Which action executes first? Action A? Action Z? action Beta?
 - Answer: We don't know
 - And we don't care! 😊
 - If we cared, the **concurrency model** was not done well
- There are points in the execution where we might care about the order of execution between tasks. These points are called **synchronization points**

Concurrency: Actions, action sequences and tasks



- The first bar in the synchronization point is called a **join** since it joins a set of actions from different tasks into a single thread
- The second bar is called a **fork** as it forks from the single thread into multiple
- A synchronization point indicates that the flow cannot proceed past the synchronization point until ALL predecessor actions have completed
- Once the three actions have completed, control goes on to subsequent actions
- Synchronization points are places in the code in which tasks will share data or ensure preconditions are met before proceeding

Concurrency: Communication between tasks

- Tasks can communicate in two fundamental ways
 - Synchronous communications
 - Asynchronous communications
- **Synchronous communications** are like phone calls
 - Both parties are engaged at the same time and invoke services and exchange data immediately
 - The sender sends and waits for a response
 - Synchronous communication is implemented via function calls

Concurrency: Communication between tasks

- Tasks can communicate in two fundamental ways
 - Synchronous communications
 - Asynchronous communications
- **Asynchronous communication** is like a postcard
 - The sender “sends and forgets”
 - At some later time, the receiver gets and processes the message or data
 - The message processing is less timely but the sender need not wait until the receiver is ready

Concurrency

- For tasks to execute concurrently, they must run on different CPUs or cores (multi-core CPU)
- Within a given CPU core, multiple tasks are run using pseudo-concurrency
 - Only one task executes at any time instant
 - The processor gives the appearance of concurrency by switching among the ready tasks
 - There are many techniques that can be used to select the task to run
 - These are known as **scheduling algorithms**

Outline

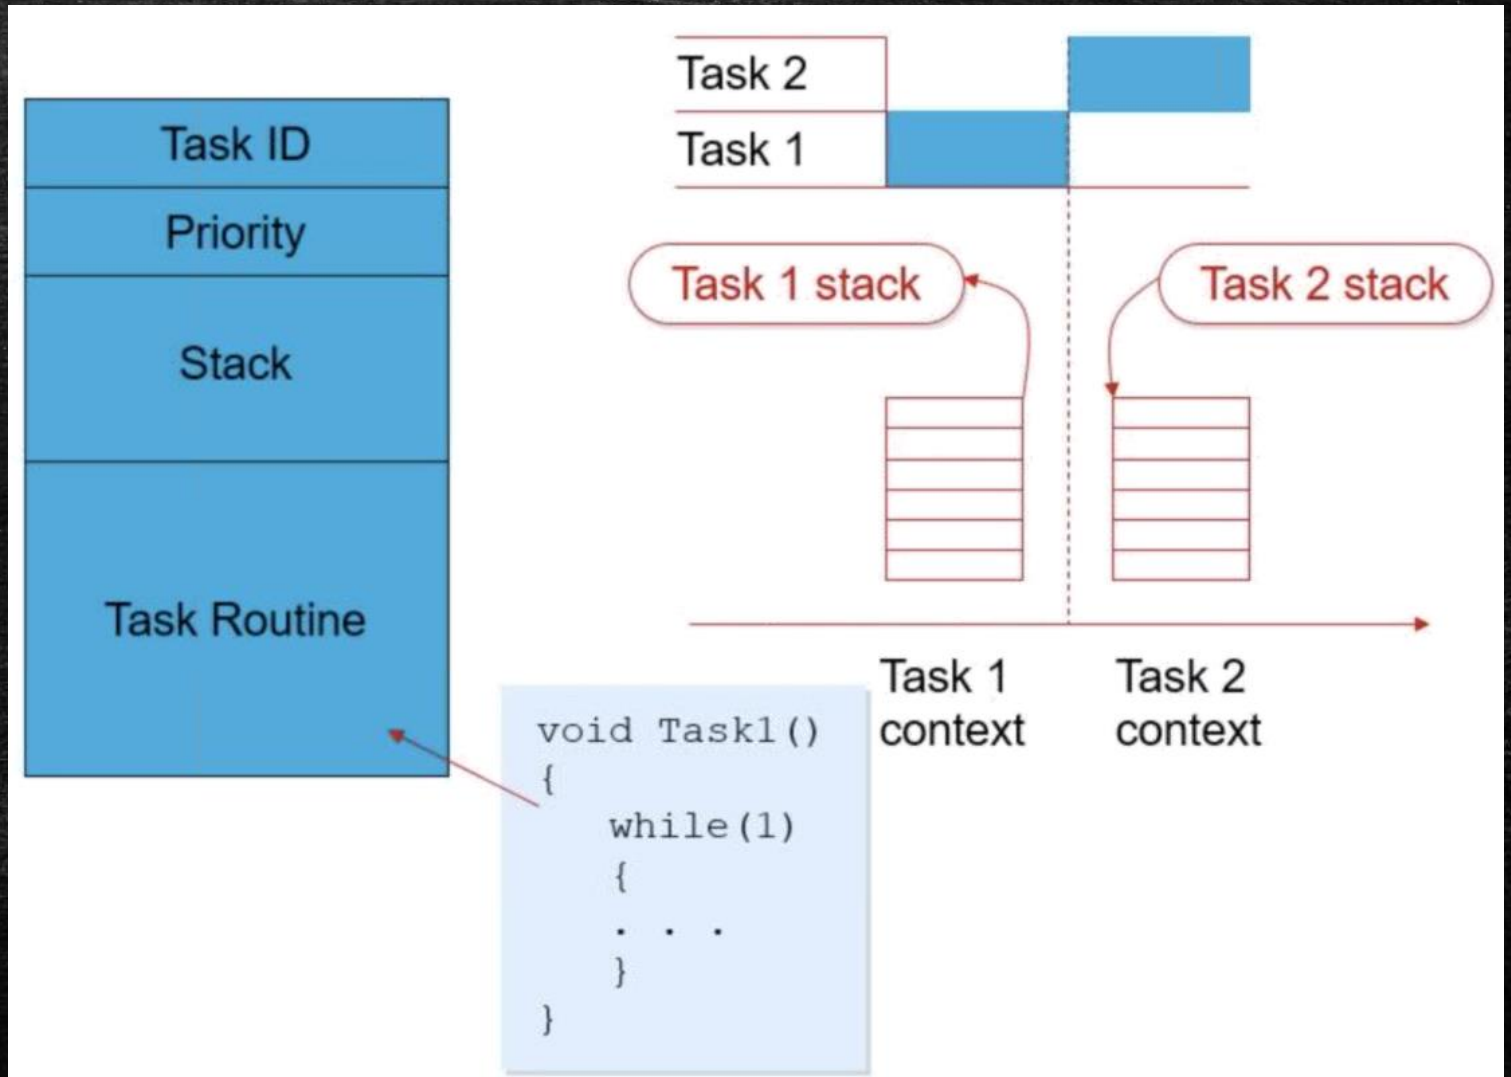
- Introduction
- Concurrency
- Real-time planning algorithms
- Reentrant code
- Design patterns

Scheduling

- Scheduling refers to the way the execution of concurrent tasks is performed
- Different techniques exist and the way they schedule tasks have pros and cons
- However, the main goal of any scheduling algorithm is to guarantee that a waiting task is given CPU time when an external event occurs
- With a single CPU, one task runs in the CPU while other tasks are waiting
- The act of reassigning a CPU from one task to another one is called **context switch**

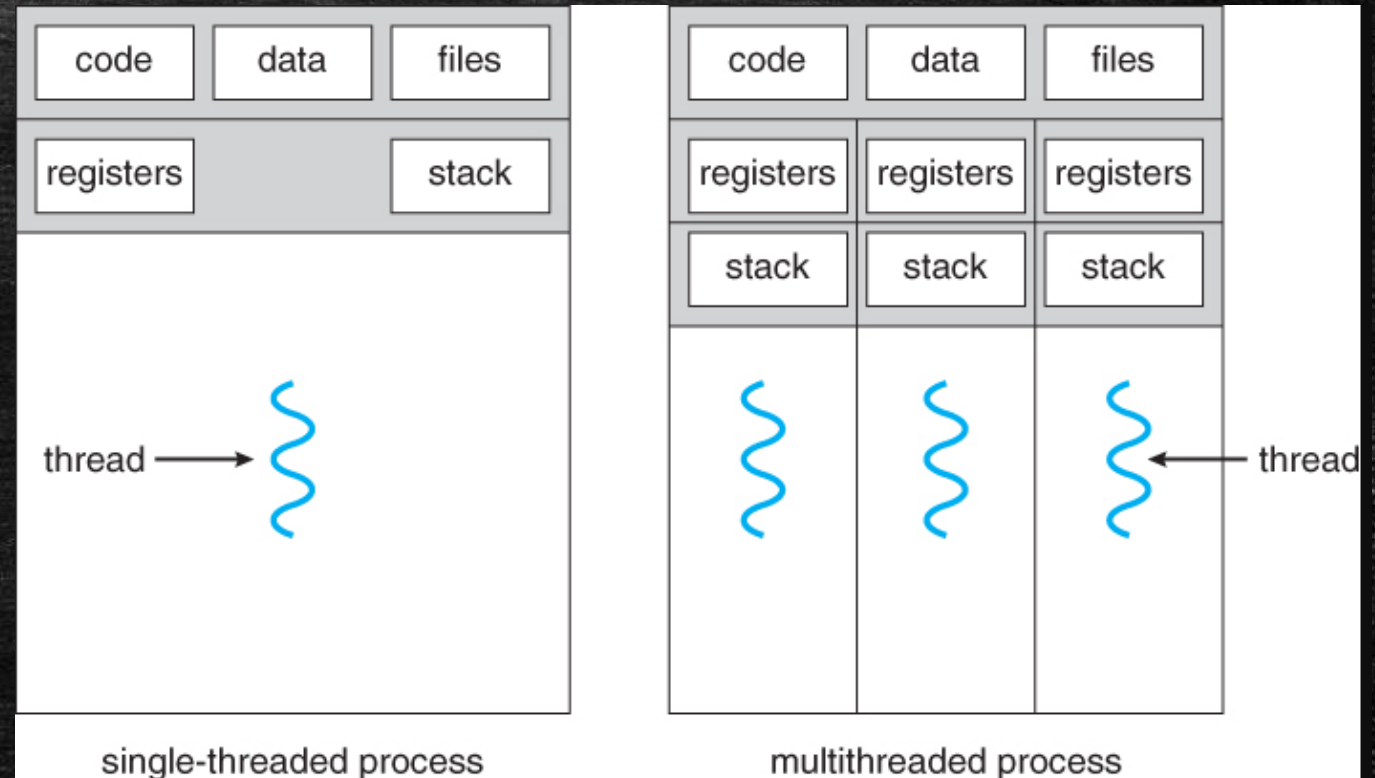
Scheduling

- Example of task switching



Scheduling

- A task represents an activity such as a **process** or a **thread**
- **Threads** are lightweight processes which share an entire memory space
- Threads are processes that run in the same memory context and may share the same data while executing
- A task can have multiple threads

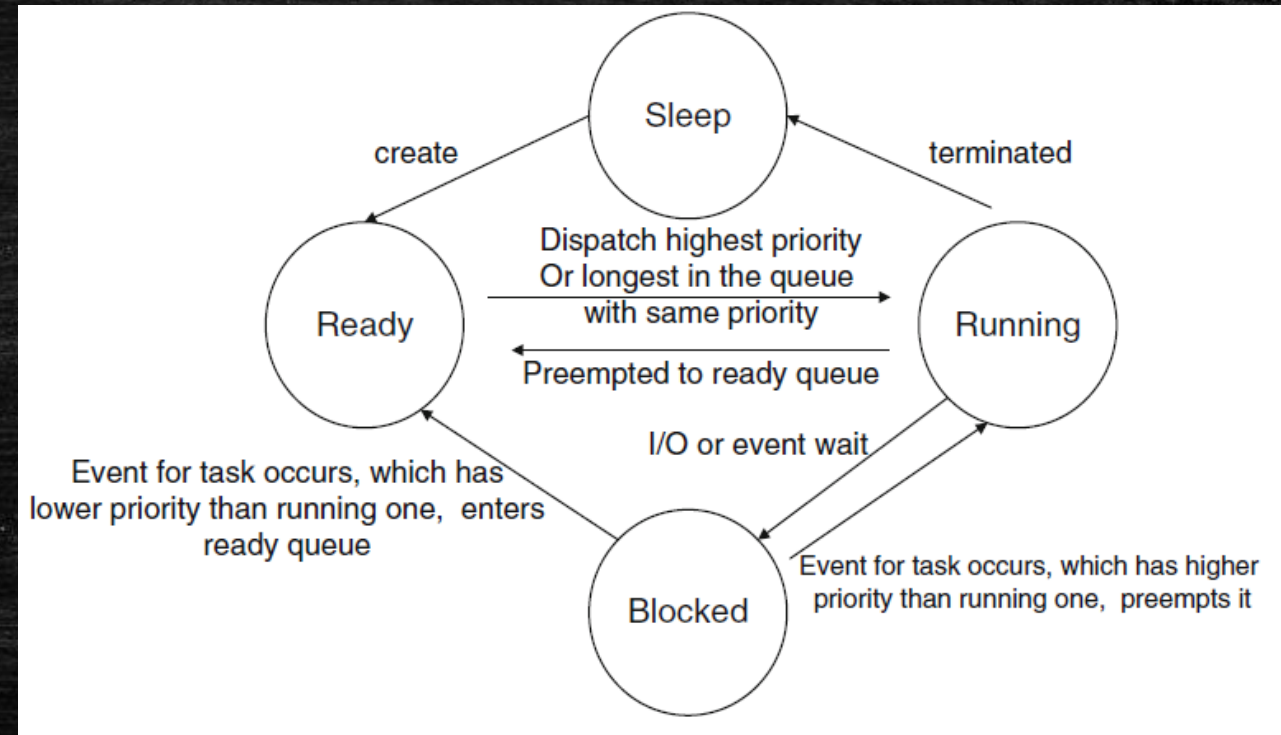
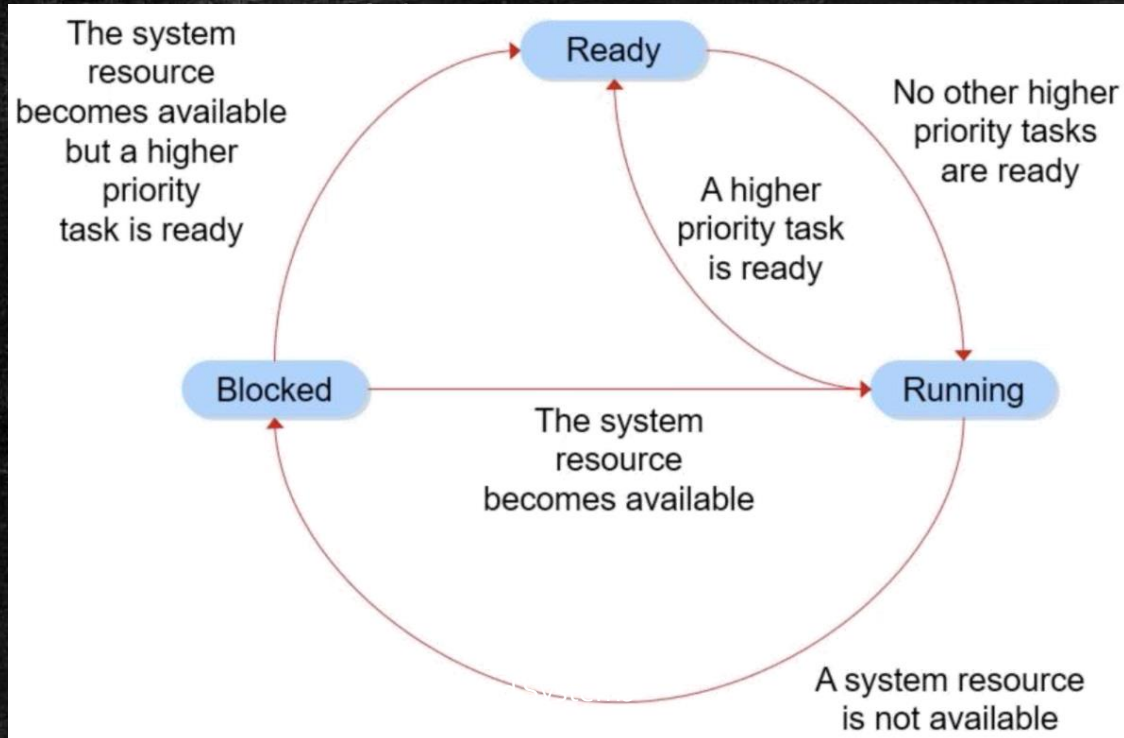


Scheduling – Timing requirement

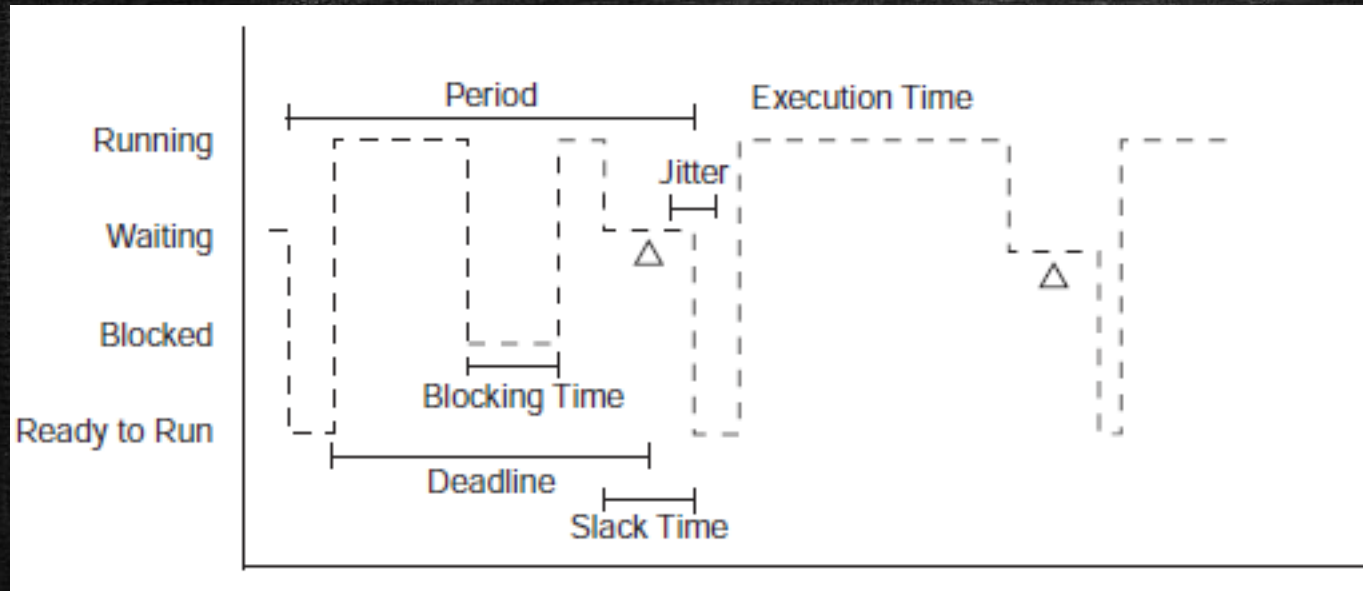
- Each task must keep track of its own concerns: CPU register status, program count, memory space and stack
- This way, the CPU can switch back and forth between these tasks
- When an event request occurs in a reactive embedded system, the target task must respond to the event within a preset time
- A parameter of a task called **Worst-Case Execution Time (WCET)** indicates the maximum length of time a task could take to execute
- A task often shows a certain variation of execution times depending on the input data or the conditions of the embedded system
- The actual response time may be shorter than the WCET

Scheduling – Task life cycle

- Every task in an embedded system has a life cycle
- Life cycle is important to determine time requirements of the task
- Depending on the scheduling algorithm/RTOS used, task life cycle



Scheduling – Time related terminology



The life cycle of a task as a function of time

- Assume a periodic task (it executes at a more or less fixed rate)
- Terminology
 - Period: duration of the time interval between task initiations
 - Jitter: variation in the period
 - Execution time: time required to complete the activities of the task
 - Deadline: period of time between when the task becomes ready to run and when it must complete

Scheduler

- The task selection for CPU (running state) is determined by a scheduler
- A scheduler has scheduling rules to select the running task
- Example of scheduling rules (RTX_51)
 - The task with the highest priority of all tasks in the READY state is executed first
 - If several tasks of the same priority are in the READY state, the task that has been ready the longest will be the next to execute
 - Exception: round-robin scheduling
- Scheduling techniques come next...

Scheduling techniques – a brief intro

- Cyclic executive scheduling

- The scheduler consists of an infinite loop that executes the tasks one after the other
- It is not flexible and can have poor response to incoming events since an event processed by a task must wait until that task is run within the cycle
- Task execution must be short if other tasks are going to run in a timely fashion

```
void main(void) {  
    /* global static and stack data */  
    static int nTasks = 3;  
    int currentTask;  
  
    /* initialization code */  
    currentTask = 0;  
  
    if (POST()) { /* Power On Self Test succeeds */  
        /* scheduling executive */  
        while (TRUE) {  
            task1();  
            task2();  
            task3();  
        }; /* end cyclic processing loop */  
    }  
}
```


Scheduling technique – a brief intro

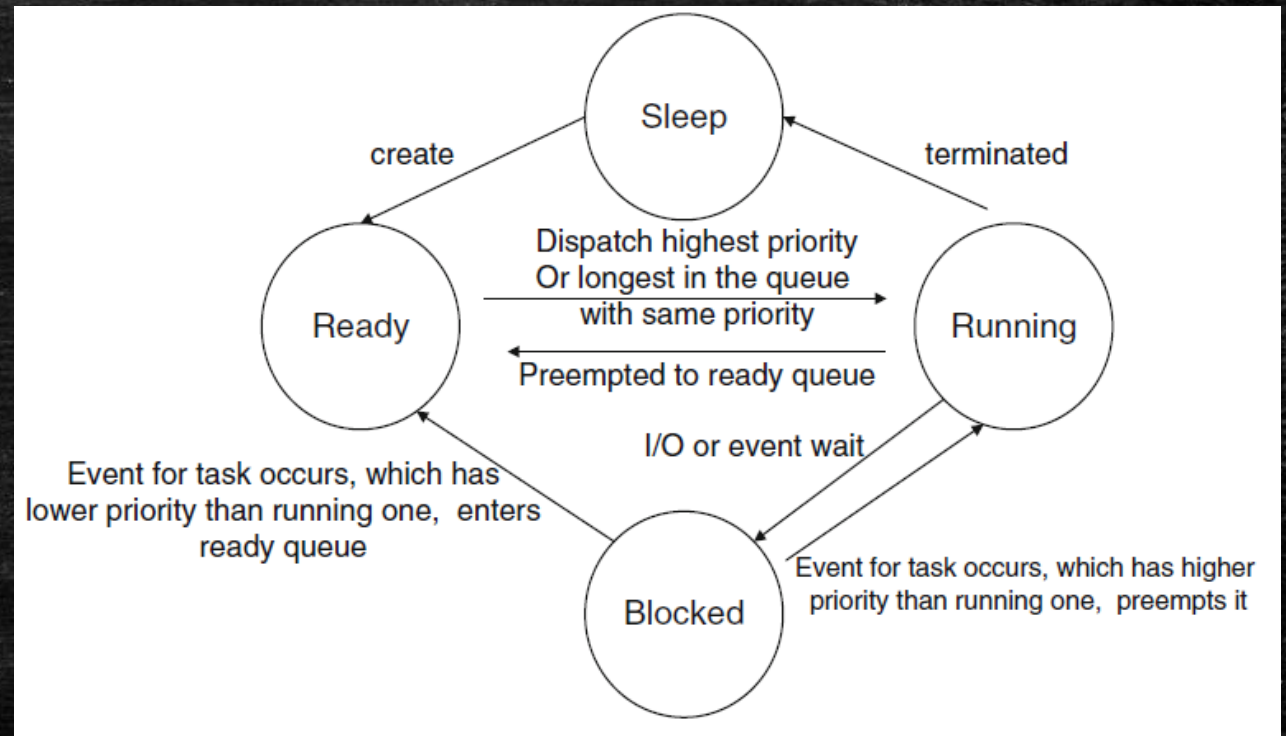
- Time-triggered cyclic executive
 - The cycle starts at a time boundary
 - If the schedule completes the cycle before the next time boundary, it sleeps until the start of the next epoch
- Cooperative round-robin
 - it is a cyclic executive in which the tasks don't run to completion but instead run to specific points at which they release control back to the scheduling loop
 - This allows long tasks to run while still permitting other tasks to make progress in their processing

Scheduling technique – a brief intro

- Cyclic and round-robin schedulers present a number of problems
- They are not responsive to urgent events; the task that processes the urgent event simply runs in its sequence
- They are not flexible since the scheduler order is determined at design time
- They don't scale well to large number of tasks
- A single misbehaving task stops the entire system from execution
- These schedulers are best suited to very simple systems

Scheduling technique – a brief intro

- The other primary kind of scheduling is **preemptive scheduling**
- The scheduler **stops** the currently executing task when it decides to do so and runs the next task
- The criterion for preempting a task and the selection of the next task in sequence depends on the particular scheduling rules
- Task life cycle presented in slide 27 are examples of preemptive scheduling




Scheduling technique – a brief intro

- The most common scheduling technique is priority based scheduling
- Each task is assigned a priority, a scalar value that will be used to select the task to execute
- The task runs in an infinite loop and will be interrupted by the scheduler when it is appropriate for other tasks to run


```
/*  
    here is where private static data goes  
*/  
static int taskAInvocationCounter = 0;  
  
void taskA(void) {  
    /* more private task data */  
    int stuffToDo;  
  
    /* initialization code */  
    stuffToDo = 1;  
  
    while (stuffToDo) {  
        signal = waitOnSignal();  
        switch (signal) {  
            case signal1:  
                /* signal 1 processing here */  
                break;  
            case signal2:  
                /* signal 2 processing here */  
            case signal3:  
                /* signal 3 processing here */  
        };  
    }; /* end infinite while loop */  
}; /* end task */
```


Scheduling Techniques – a brief intro

Round Robin



Priority-based



Non-Preemptive Scheduling

- Simple task scheduling method for periodic time requirement systems
- Systems belonging to this type of scheduling are periodic systems or systems where “cooperative” multitasking is implemented
- In **cooperative multitasking** the process that is running must offer control to other processes/tasks after a given amount of time
- Another example is cyclic scheduling where tasks can be scheduled in a fixed static table

Non-Preemptive Scheduling

- Non-preemptive scheduling is typically used in embedded systems that have a set of tasks and all of them have a fixed period.
- Aperiodic events can be estimated by their worst case interval gap between two consecutive task events
- All tasks are independent and WCET is known in advance, so that the deadline is at the end of the WCET
- The advantage of non-preemptive scheduling is zero-overhead for context switches between process/tasks
- The disadvantage is the inflexibility

Non-Preemptive Scheduling

- If there are n tasks and the WCET of the i^{th} task is c_i , then the period of any task T_i must satisfy

$$T_i \geq \sum c_j \quad (j = 1, 2, \dots, n)$$

- Otherwise some tasks may miss its deadline
- As a result, software design should make each task as shorter as possible. For long tasks, they should be broken into many small tasks

Non-Preemptive Scheduling

- Cyclic schedule example

Task	Execution time	Period
Task 1	20 ms	50 ms
Task 2	25 ms	100 ms

- Period of T₁ (50ms) \geq c₁ (20ms) + c₂(25ms)
- Period of T₂ (100ms) \geq c₁ (20ms) + c₂(25ms)
- The tasks are **schedulable**, they wont miss their deadline


Non-Preemptive Scheduling

- Cyclic schedule example
 - Assumptions: there is a timer set on the period of 50 ms

```
while (1)
{
    Wait_for_50ms_timer_interrupt;
    do task1;
    Wait_for_50ms_timer_interrupt;
    do task1;
    do task2;
}
```


Pre-emptive Scheduling

- Priority-driven scheduling
- Reactive embedded systems must respond and handle external or internal events with different urgency
- Some events have hard real-time constraints while the others may only have soft real-time constraints
- That is, some tasks should be assigned higher priority than other tasks



Pre-emptive Scheduling

- The CPU always goes to the highest priority process that is ready
- Pre-emptive/priority-based scheduling has several algorithms
 - Static priority-based
 - Static timing scheduling
 - Round-robin scheduling
 - Rate Monotonic Scheduling (RMS) – based on the deadline based analysis called Rate Monotonic Analysis (RMA)
 - Dynamic priority-based
 - Earliest Deadline First (EDF): it assigns priorities at runtime based on upcoming execution deadlines
 - Priority is time varying rather than fixed

Rate Monotonic Scheduling

- Priority-based static scheduling method for preemptive real-time systems
- It can guarantee the time requirement and maximize the schedulability as long as the CPU utilization is below 0.7
- RMA (Rate Monotonic Analysis) has proven to be optimal among all static priority scheduling algorithms

Rate Monotonic Scheduling

- It assigns shorter period/deadline processes higher priority at design time, assuming that the processes have the following properties
 - No resource sharing
 - Deterministic deadlines are exactly equal to periods
 - Context switch times are free and have no impact on the model
- Once the priority of a task is assigned, it will remain constant for the lifetime of the task

Rate Monotonic Scheduling

- Example 1

Task	Execution Time	Period (Deadline)	Priority	Utilization
Task 1	20 ms	50 ms	2 (high)	40%
Task 2	45 ms	100 ms	1 (low)	45%

- Task 1 must meet its deadline of 50, 100, 150, ...
- Task 2 must meet its deadline at 100, 200, 300, ...
- Task 1 is assigned higher priority cause it has shorter period

Rate Monotonic Scheduling

- Example

Task	Execution Time	Period (Deadline)	Priority	Utilization
Task 1	20 ms	50 ms	2 (high)	40%
Task 2	45 ms	100 ms	1 (low)	45%

- Schedule:

T1(20) T2(30) T1(20) T2(25) (5) T1(20) T2(30) Continue the same pattern

50

100

150

Rate Monotonic Scheduling

- Schedule:

T1(20) T2(30) T1(20) T2(15) (15) T1(20) T2(30) Continue the same pattern

50

100

150

- At time 50, task 2 is preempted by task 1 cause it is ready every 50ms and has higher priority
- This schedule guarantees all tasks meet their deadline

Rate Monotonic Scheduling

- If task 2 gets higher priority over task 1, then task 1 will miss its deadline at time 50

T2(45) T1(5) T1(45)

50

Rate Monotonic Scheduling

- Example 2

Task	Execution time	Period (Deadline)	Priority	Utilization
Task 1	4 ms	10 ms	3 (high)	40%
Task 2	5 ms	15 ms	2 (medium)	33%
Task 3	6 ms	25 ms	1 (low)	24%

- Periods of tasks obey $T(1) < T(2) < T(3)$
- Therefore $P(1) > P(2) > P(3)$
- CPU utilization is 97%

Rate Monotonic Scheduling

- Schedule:

T ₁ (4)	T ₂ (5)	T ₃ (1)	T ₁ (4)	T ₃ (1)	T ₂ (5)	T ₁ (4)	T ₃ (1)	-3
		10		15	20		25	


- At time 25, task 3 misses its deadline by 3 ms
- This is due to the fact that total utilization is 97% which is beyond the schedulable bound of 70%

Rate Monotonic Scheduling

- Example 2: Let's change the utilization rate

Task	Execution time	Period (Deadline)	Priority	Utilization
Task 1	3 ms	10 ms	3 (high)	30%
Task 2	4 ms	15 ms	2 (medium)	26%
Task 3	4 ms	25 ms	1 (low)	16%

- CPU utilization is 72%



It is not always possible to fully utilize the CPU

Rate Monotonic Scheduling

- According to Rate Monotonic Analysis, fixed-priority scheduling has a **worst case schedulable** bound of

$$U_n = \sum_{i=1}^n \frac{c_i}{T_i} \leq n^{2^{1/n}-1}$$

n is the number of tasks in the system

- As the number of tasks increases, the schedulable bound decreases, eventually approaching its limit of $\ln 2 = 69.3\%$
- It has been shown that for a set of n periodic tasks with unique periods, a feasible schedule that will meet deadlines exist if **CPU utilization is $< U_n$**

Rate Monotonic Scheduling

- The other 30% of CPU utilization can be dedicated to lower-priority non real-time tasks
- The context switch cost for the RMS is very high
- In order to fully make use of CPU time and also to meet all deadlines, we need to use a **dynamic priority scheduling algorithm**

Dynamic Scheduling with EDF

- In a dynamic priority-based scheduling the priority of a process changes in order to increase the CPU utilization and to allow all tasks meet their deadline
- In Earliest Deadline First (EDF), higher priority is assigned to those tasks that are closer to their deadline at runtime
- EDF can be applied to both periodic and aperiodic tasks if deadlines are known in advance
- EDF is not very practical and is not as popular as RMS

Dynamic Scheduling with EDF

- In a dynamic priority-based scheduling the priority of a process changes in order to increase the CPU utilization and to allow all tasks meet their deadline
- In Earliest Deadline First (EDF), higher priority is assigned to those tasks that are closer to their deadline at runtime
- EDF can be applied to both periodic and aperiodic tasks if deadlines are known in advance
- EDF is not very practical and is not as popular as RMS
- EDF must recalculate the priority of each process at every context switch time (preemption time)

RTOS revisited
